

# Regularization and logistic regression

## BitTiger DS501 Week 3 HW Meina Wang

### Question 1

Suppose we fit “Lasso Regression” to a data set, which has 100 features ( $X_1, X_2 \dots X_{100}$ ). Now, we rescale one of these feature by multiplying with 10 (say that feature is  $X_1$ ), and then refit Lasso regression with the same regularization parameter. Now, which of the following option will be correct?

- A. It is more likely for  $X_1$  to be excluded from the model
- B. It is more likely for  $X_1$  to be included in the model
- C. Can't say
- D. None of these

### Answer 1

B.

Multiplying the feature by 10 will lead to smaller coefficient assigned to this feature. Smaller coef will have less penalty by LASSO, thus would be more likely to be included in the model.

### Question 2

Suppose you have fitted a multiple regression model on a dataset. Now, you are using Ridge regression with tuning parameter  $\lambda$  to reduce its complexity. Choose the options below which describes relationship of bias and variance with  $\lambda$ .

- A. In case of very large  $\lambda$ ; bias is low, variance is low
- B. In case of very large  $\lambda$ ; bias is low, variance is high
- C. In case of very large  $\lambda$ ; bias is high, variance is low
- D. In case of very large  $\lambda$ ; bias is high, variance is high

### Answer 2

C.

Very large  $\lambda$  would penalize the parameters heavily, thus decreasing variance, but increasing bias (likely to underfit)

## Question 3

Write a function to realize gradient descent in R. Understand how learning rate affects convergence.

## Answer 3

From class slides, gradient descent steps:

- Step 0: find an initial  $\beta_j^{(0)}$
- Step 1:  $\beta_j^{(t+1)} := \beta_j^{(t)} - \eta \frac{\partial l(\beta)}{\partial \beta_j}$ , where  $\eta$  is learning rate, and  $\frac{\partial l(\beta)}{\partial \beta_j}$  is gradient.
- step 2: check if  $\Delta_\beta l(\beta) \leq \epsilon$ , where  $\epsilon$  is the convergence threshold. If not, repeat step 1.

In [2]:

```
gradientDescent<-function(y, X, learning_rate, convergence_threshold, iters){

  X = as.matrix(data.frame(rep(1,length(y)),X))
  N = dim(X)[1]
  beta.init = as.matrix(rnorm(n=dim(X)[2], mean=0,sd = 1)) # initialize beta
  beta.init = t(beta.init)
  error = t(y) - beta.init%*%t(X)
  grad.init = -(2/N)%*%(error)%*%X # initialize gradient # %*% is operator for the
  beta = beta.init - learning_rate*(1/N)*grad.init
  convergence_threshold = 0.0001
  for(i in 1:iters){ # starts from 1
    error = t(y) - beta%*%t(X)
    grad = -(2/N)%*%error%*%X
    beta = tbeta - learning_rate*(2/N)*grad
    if(sqrt(sum(grad^2)) <= convergence_threshold){
      break
    }
  }
  print("Algorithm converged")
  print(paste("Final gradient norm is",sqrt(sum(grad^2))))
  values<-list("coef" = t(beta), "l2loss" = l2loss)
  return(values)
}
```

Learning rate determines how fast or slow each step takes towards the optimal weights (minimize cost function). If learning rate is very large we will skip the optimal solution (global minimum). If it is too small we might be trapped at local minimum, or will need too many iterations to converge to the best values.

## Question 4

A five year follow-up study on 600 disease free subjects was carried out to assess the effect of whether having exposure E or not (of smoking for example) on the development (or not) of a certain disease. The variables AGE (continuous) and obesity status (boolean), which were determined at the start of the follow-up and were to be considered as control variables in analyzing the data.

(1) State the logit form of a logistic regression model that assesses the effect of the 0/1 exposure variable E controlling for the confounding effects of AGE and OBS and the interaction effects of AGE with E and OBS with E.

(2) Given above model you have, give a formula for the odds ratio for the exposure-disease relationship that controls for the confounding and interactive effects of AGE and OBS.

(3) Now use the formula from above to write an expression for the estimated odds ratio for the exposure-disease relationship when AGE=40 and OBS=1.

## Answer 4

(1)

$$\text{logit}(P) = \beta_0 + \beta_1 * E + \beta_2 * AGE + \beta_3 * OBS + \beta_4 * AGE * E + \beta_5 * OBS * E$$

P is the probability of the development of a certain disease.

E is the exposure variable, and AGE and OBS are the control variables.  $\beta$  is the coefficient for each variable.

The last two terms on the right hand side of the equation above are interaction effects of AGE with E and OBS with E.

(2)

$$\text{Odds ratio for the exposure-disease relationship} = \frac{p}{1-p} = \frac{\text{probability of disease WITH exposure}}{\text{probability of disease WITHOUT exposure}}$$

$$\text{logit}(P) = \log\left(\frac{p}{1-p}\right) = \log(\text{odds})$$

Taking e to the power to both sides of the above equation yields,

$$e^{\log(\text{odds})} = \text{odds} = e^{\log\left(\frac{p}{1-p}\right)}$$

Thus, the odds ratio for comparing WITH and WITHOUT exposure is

$$\begin{aligned} \frac{\text{odds}_{\text{WITH exposure}}}{\text{odds}_{\text{WITHOUT exposure}}} &= \frac{\frac{p_1}{1-p_1}}{\frac{p_2}{1-p_2}} = \frac{\exp(\text{logit}(P_1))}{\exp(\text{logit}(P_2))} = \exp(\text{logit}(P_1) - \text{logit}(P_2)) = \\ &\exp((\beta_0 + \beta_1 * 1 + \beta_2 * AGE_1 + \beta_3 * OBS_1 + \beta_4 * AGE_1 * 1 + \beta_5 * OBS_1 * 1) - (\beta_0 + \beta_1 * 0 + \beta_2 * AGE_2 + \beta_3 * OBS_2 + \beta_4 * AGE_2 * 0 + \beta_5 * OBS_2 * 0)) = \\ &\exp((\beta_0 + \beta_1 + \beta_2 * AGE_1 + \beta_3 * OBS_1 + \beta_4 * AGE_1 + \beta_5 * OBS_1) - (\beta_0 + \beta_2 * AGE_2 + \beta_3 * OBS_2)) = \\ &\exp(\beta_1 + \beta_2 * (AGE_1 - AGE_2) + \beta_3 * (OBS_1 - OBS_2) + \beta_4 * AGE_1 + \beta_5 * OBS_1) \end{aligned}$$

Therefore, the odds ratio is

$\exp(\beta_1 + \beta_2 * (AGE_1 - AGE_2) + \beta_3 * (OBS_1 - OBS_2) + \beta_4 * AGE_1 + \beta_5 * OBS_1)$ , where subscript 1 indicates the AGE and OBS values when E is WITH exposure, while subscript 2 indicates the values when E is WITHOUT exposure.

(3)

Plug in AGE = 40 and OBS = 1 into the equation from (2) yields,

$$\exp(\beta_1 + \beta_2 * (40 - 40) + \beta_3 * (1 - 1) + \beta_4 * 40 + \beta_5 * 1) = \exp(\beta_1 + \beta_4 * 40 + \beta_5 * 1)$$

## Question 5

Build the best logistic regression model to predict loan will be default (delay) or not. Add regularization to control for multicollinearity.

## Answer 5

In [1]:

```
# load the data
loan <- read.csv("/users/meinawang/Documents/bittiger/DS501/lending-club-loan-data/")
loanT <- loan
```

In [2]:

```
# select only features with no more than 20% missing data
loan$dti <- ifelse(!is.na(loan$dti_joint), loan$dti_joint, loan$dti)
loan$annual_inc <- ifelse(!is.na(loan$annual_inc_joint), loan$annual_inc_joint, loan$annual_inc)
num.NA <- sort(sapply(loan, function(x) {sum(is.na(x))}), decreasing=TRUE)
remain.col <- names(num.NA)[which(num.NA <= 0.8 * dim(loan)[1])]
loan <- loan[, remain.col]
```

```
In [3]:  
  
# based on last week's hw, log transform features so that they are close to normal  
loan$home_ownership <- ifelse(loan$home_ownership %in% c('ANY', 'NONE', 'OTHER'), 'OTHER',  
                             loan$home_ownership)  
int_state <- by(loan, loan$addr_state, function(x) {  
  return(mean(x$int_rate))  
})  
loan$state_mean_int <-  
  ifelse(loan$addr_state %in% names(int_state)[which(int_state <= quantile(int_state, 0.25))],  
         ifelse(loan$addr_state %in% names(int_state)[which(int_state <= quantile(int_state, 0.5))],  
               ifelse(loan$addr_state %in% names(int_state)[which(int_state <= quantile(int_state, 0.75))],  
                     'mediumhigh', 'high'))  
loan$tot_cur_bal[which(is.na(loan$tot_cur_bal))] <- median(loan$tot_cur_bal, na.rm = T)  
loan$total_acc[which(is.na(loan$total_acc))] <- median(loan$total_acc, na.rm = T)  
loan$open_acc[which(is.na(loan$open_acc))] <- median(loan$open_acc, na.rm = T)  
loan$annual_inc[which(is.na(loan$annual_inc))] <- median(loan$annual_inc, na.rm = T)  
loan$tot_coll_amt[which(is.na(loan$tot_coll_amt))] <- median(loan$tot_coll_amt, na.rm = T)  
loan$inq_last_6mths[which(is.na(loan$inq_last_6mths))] <- median(loan$inq_last_6mths, na.rm = T)  
loan$log_annual_inc <- log(loan$annual_inc + 1)  
loan$log_last_pymnt_amnt <- log(loan$last_pymnt_amnt + 1)  
loan$log_tot_coll_amt <- log(loan$tot_coll_amt + 1)  
loan$total_rev_hi_lim_log = log(loan$total_rev_hi_lim + 1)  
loan$total_rec_int_log = log(loan$total_rec_int + 1)
```

```
In [4]:  
  
# check the levels from feature loan_status  
loan$loan_status <- gsub('Does not meet the credit policy. Status:', '', loan$loan_status)  
sort(table(loan$loan_status))  
  
nums <- sapply(loan, is.numeric)  
loan_nums <- loan[, nums]  
  
categories <- sapply(loan, is.character)  
loan_categories <- loan[, categories]
```

	Default	Late (16-30 days)	In Grace Period	
Issued	1219	2357	6253	
8460				
Late (31-120 days)		Charged Off	Fully Paid	Cu
Current	11591	46009	209711	6
01779				

In [6]:

```
# Convert the levels in loan_status into binary, 'Default', 'Charged Off', 'Late (1-15 days)',
#'Late (31-120 days)', 'Charged Off' are 1, the rest levels are 0.

loan$loan_status_binary <- with(loan, ifelse(loan_status %in% c('Default', 'Charged Off', 'Late (1-15 days)',
                                                                'In Grace Period', 'Charged Off'), 1, 0))

table(loan$loan_status_binary)

loan$loan_status_binary = factor(loan$loan_status_binary)
levels(loan$loan_status_binary)
```

```
      0      1
819950 67429
```

```
'0' '1'
```


In [20]:

```
# check the pdf of features, and see if they are normal.
library(repr)

# Change plot size to 6 x 3
options(repr.plot.width=6, repr.plot.height=3)
```

In [22]:

```
# check how the numerical features response to loan_status_binary  
boxplot(total_acc ~ loan_status_binary, data = loan,xlab='loan_status_binary',ylab='total_acc')  
boxplot(log_tot_coll_amt ~ loan_status_binary, data = loan,xlab='loan_status_binary',ylab='log_tot_coll_amt')  
boxplot(tot_cur_bal ~ loan_status_binary, data = loan,xlab='loan_status_binary',ylab='tot_cur_bal')  
boxplot(inq_last_6mths ~ loan_status_binary, data = loan,xlab='loan_status_binary',ylab='inq_last_6mths')  
boxplot(open_acc ~ loan_status_binary, data = loan,xlab='loan_status_binary',ylab='open_acc')  
boxplot(log_annual_inc ~ loan_status_binary, data = loan,xlab='loan_status_binary',ylab='log_annual_inc')  
boxplot(total_pymnt ~ loan_status_binary, data = loan,xlab='loan_status_binary',ylab='total_pymnt')  
boxplot(total_rec_int_log ~ loan_status_binary, data = loan,xlab='loan_status_binary',ylab='total_rec_int_log')  
boxplot(log_last_pymnt_amnt ~ loan_status_binary, data = loan,xlab='loan_status_binary',ylab='log_last_pymnt_amnt')  
boxplot(int_rate ~ loan_status_binary, data = loan,xlab='loan_status_binary',ylab='int_rate')  
boxplot(total_rec_int_log ~ loan_status_binary, data = loan,xlab='loan_status_binary',ylab='total_rec_int_log')
```



In [12]:

```
library(MASS)
# compute the chi square value between categorical features with loan_status_binary
chisq.test(table(loan$term, loan$loan_status_binary))
chisq.test(table(loan$home_ownership, loan$loan_status_binary))
chisq.test(table(loan$application_type, loan$loan_status_binary))
chisq.test(table(loan$pymnt_plan, loan$loan_status_binary))
chisq.test(table(loan$state_mean_int, loan$loan_status_binary))
chisq.test(table(loan$emp_length, loan$loan_status_binary))
```

Pearson's Chi-squared test with Yates' continuity correction

```
data: table(loan$term, loan$loan_status_binary)
X-squared = 1390.6, df = 1, p-value < 2.2e-16
```

Pearson's Chi-squared test

```
data: table(loan$home_ownership, loan$loan_status_binary)
X-squared = 1333.6, df = 3, p-value < 2.2e-16
```

Pearson's Chi-squared test with Yates' continuity correction

```
data: table(loan$application_type, loan$loan_status_binary)
X-squared = 30.979, df = 1, p-value = 2.609e-08
```

```
Warning message in chisq.test(table(loan$pymnt_plan, loan$loan_status_
binary)):
"Chi-squared approximation may be incorrect"
```

Pearson's Chi-squared test with Yates' continuity correction

```
data: table(loan$pymnt_plan, loan$loan_status_binary)
X-squared = 19.923, df = 1, p-value = 8.061e-06
```

Pearson's Chi-squared test

```
data: table(loan$state_mean_int, loan$loan_status_binary)
X-squared = 150.14, df = 3, p-value < 2.2e-16
```

Pearson's Chi-squared test

```
data: table(loan$emp_length, loan$loan_status_binary)
X-squared = 512.67, df = 11, p-value < 2.2e-16
```



In [13]:

```
library(glmnet) # can only take matrix
```

Loading required package: Matrix

Loading required package: foreach

Warning message:

"package 'foreach' was built under R version 3.4.3"Loaded glmnet 2.0-1  
3

In [14]:

```
# scale (normalize) all the numerical features before doing train / test split.
```

```
nums <- sapply(loan, is.numeric)
```

```
loan[nums] <- lapply(loan[nums], scale)
```

```
summary(loan)
```

mths_since_last_major_derog	mths_since_last_delinq	tot_coll_amt	
Min. : 0.0	Min. : 0.0	Min. : 0	
1st Qu.: 27.0	1st Qu.: 15.0	1st Qu.: 0	
Median : 44.0	Median : 31.0	Median : 0	
Mean : 44.1	Mean : 34.1	Mean : 208	
3rd Qu.: 61.0	3rd Qu.: 50.0	3rd Qu.: 0	
Max. :188.0	Max. :188.0	Max. :9152545	
NA's :665676	NA's :454312		
tot_cur_bal	total_rev_hi_lim	revol_util	
Min. : 0	Min. : 0	Min. : 0.00	
1st Qu.: 32246	1st Qu.: 13900	1st Qu.: 37.70	
Median : 80559	Median : 23700	Median : 56.00	
Mean : 134794	Mean : 32069	Mean : 55.07	
3rd Qu.: 195794	3rd Qu.: 39800	3rd Qu.: 73.60	
Max. :8000078	Max. :9999999	Max. :892.30	
	NA's :70276	NA's :502	
collections_12_mths_ex_med	delinq_2yrs	inq_last_6mths	open_acc
Min. : 0.00000	Min. : 0.0000	Min. : 0.0000	Min. :

In [15]:

```
set.seed(1)
```

```
train.ind <- sample(1:dim(loan)[1], 0.7 * dim(loan)[1])
```

```
train <- loan[train.ind, ]
```

```
dim(train)
```

```
test <- loan[-train.ind, ]
```

In [16]:

```
# select the numerical/categorical features that are responsive to loan_status_bina
train.sub <- train[, c('loan_status_binary', 'total_acc', 'tot_cur_bal', 'inq_last_6mths',
                      'int_rate', 'total_pymnt', 'total_rec_int_log', 'log_last_pymnt',
                      'open_acc', 'log_tot_coll_amt', 'term', 'home_ownership', 'app',
                      'state_mean_int', 'emp_length')]
```

```
test.sub <- test[, c('loan_status_binary', 'total_acc', 'tot_cur_bal', 'inq_last_6mths',
                    'int_rate', 'total_pymnt', 'total_rec_int_log', 'log_last_pymnt',
                    'open_acc', 'log_tot_coll_amt', 'term', 'home_ownership', 'app',
                    'state_mean_int', 'emp_length')]
```

In [17]:

```
# convert to matrixes as required by glmnet.
ind <- train.sub[, -1]
ind <- model.matrix(~., ind)

dim(ind)
summary(ind)
dep <- train.sub[, 1]
dep <- as.matrix(dep)
summary(dep)
dim(dep)
typeof(dep)
```

621165 32

(Intercept)		total_acc	tot_cur_bal	inq_last_6mths
Min.	:1	Min. :-2.049597	Min. :-0.908367	Min. :-0.69569
1st Qu.:	:1	1st Qu.: -0.698288	1st Qu.: -0.690787	1st Qu.: -0.69569
Median	:1	Median :-0.107090	Median :-0.365485	Median :-0.69569
Mean	:1	Mean : 0.001167	Mean : 0.000883	Mean :-0.00137
3rd Qu.:	:1	3rd Qu.: 0.568565	3rd Qu.: 0.413477	3rd Qu.: 0.30588
Max.	:1	Max. :11.041210	Max. :31.253584	Max. :30.35277
loan_amnt		dti	int_rate	
Min.	:-1.6899223	Min. :-2.1844292	Min.	:-1.808987
1st Qu.:	-0.8008180	1st Qu.: -0.7503527	1st Qu.:	-0.743231
Median	:-0.2080818	Median :-0.0574165	Median	:-0.058591
Mean	: 0.0000272	Mean :-0.0004977	Mean	:-0.000683
3rd Qu.:	0.6217489	3rd Qu.: 0.6993902	3rd Qu.:	0.673973
Max.	: 2.3999576	Max. : 3.1011671	Max.	: 3.592820
total_pymnt		total_rec_int_log	log_last_pymnt_amnt	
Min.	:-0.960309	Min. :-4.409052	Min.	:-3.890081
1st Qu.:	-0.716842	1st Qu.: -0.432287	1st Qu.:	-0.429845
Median	:-0.338710	Median : 0.146676	Median	:-0.122226
Mean	:-0.000091	Mean : 0.000262	Mean	: 0.000569
3rd Qu.:	0.389095	3rd Qu.: 0.625063	3rd Qu.:	0.235473
Max.	: 6.380028	Max. : 2.158997	Max.	: 2.555184
log_annual_inc		open_acc	log_tot_coll_amt	term 60 months

Min. : -6.677380	Min. : -2.17189	Min. : -0.375950	Min. : 0.
1st Qu.: -0.683414	1st Qu.: -0.66735	1st Qu.: -0.375950	1st Qu.: 0.
Median : 0.012670	Median : -0.10315	Median : -0.375950	Median : 0.
Mean : 0.000566	Mean : 0.00057	Mean : 0.000384	Mean : 0.
3rd Qu.: 0.628682	3rd Qu.: 0.46106	3rd Qu.: -0.375950	3rd Qu.: 1.
Max. : 9.448524	Max. : 13.24968	Max. : 5.736659	Max. : 1.
home_ownershipOTHER	home_ownershipOWN	home_ownershipRENT	application_typeJOINT
Min. : 0.0000000	Min. : 0.00000	Min. : 0.0000	Min. : 0.00
1st Qu.: 0.0000000	1st Qu.: 0.00000	1st Qu.: 0.0000	1st Qu.: 0.00
Median : 0.0000000	Median : 0.00000	Median : 0.0000	Median : 0.00
Mean : 0.0002544	Mean : 0.09876	Mean : 0.4009	Mean : 0.00
3rd Qu.: 0.0000000	3rd Qu.: 0.00000	3rd Qu.: 1.0000	3rd Qu.: 0.00
Max. : 1.0000000	Max. : 1.00000	Max. : 1.0000	Max. : 1.00
state_mean_intlow	state_mean_intlowmedium	state_mean_intmediumhigh	
Min. : 0.00000	Min. : 0.0000	Min. : 0.0000	
1st Qu.: 0.00000	1st Qu.: 0.0000	1st Qu.: 0.0000	
Median : 0.00000	Median : 1.0000	Median : 0.0000	
Mean : 0.08751	Mean : 0.5175	Mean : 0.2654	
3rd Qu.: 0.00000	3rd Qu.: 1.0000	3rd Qu.: 1.0000	
Max. : 1.00000	Max. : 1.0000	Max. : 1.0000	
emp_length1 year	emp_length10+ years	emp_length2 years	emp_length3 years
Min. : 0.00000	Min. : 0.0000	Min. : 0.00000	Min. : 0.000
1st Qu.: 0.00000	1st Qu.: 0.0000	1st Qu.: 0.00000	1st Qu.: 0.000
Median : 0.00000	Median : 0.0000	Median : 0.00000	Median : 0.000
Mean : 0.06412	Mean : 0.3289	Mean : 0.08876	Mean : 0.078
3rd Qu.: 0.00000	3rd Qu.: 1.0000	3rd Qu.: 0.00000	3rd Qu.: 0.000
Max. : 1.00000	Max. : 1.0000	Max. : 1.00000	Max. : 1.000
emp_length4 years	emp_length5 years	emp_length6 years	emp_length7 years
Min. : 0.00000	Min. : 0.0000	Min. : 0.00000	Min. : 0.00000
1st Qu.: 0.00000	1st Qu.: 0.0000	1st Qu.: 0.00000	1st Qu.: 0.00000
Median : 0.00000	Median : 0.0000	Median : 0.00000	Median : 0.00000
Mean : 0.05882	Mean : 0.0627	Mean : 0.04861	Mean : 0.05035

3rd Qu.:0.00000	3rd Qu.:0.0000	3rd Qu.:0.00000	3rd Qu.:0.00000
Max. :1.00000	Max. :1.0000	Max. :1.00000	Max. :1.00000
emp_length8 years	emp_length9 years	emp_lengthn/a	
Min. :0.00000	Min. :0.00000	Min. :0.00000	
1st Qu.:0.00000	1st Qu.:0.00000	1st Qu.:0.00000	
Median :0.00000	Median :0.00000	Median :0.00000	
Mean :0.04962	Mean :0.03891	Mean :0.05051	
3rd Qu.:0.00000	3rd Qu.:0.00000	3rd Qu.:0.00000	
Max. :1.00000	Max. :1.00000	Max. :1.00000	

```
V1
0:574004
1: 47161
```

```
621165  1
```

```
'character'
```

```
In [18]:
```

```
# cross-validation with glmnet for logistic regression
cvfit <- cv.glmnet(ind, dep, family = "binomial", type.measure = "mse")
plot(cvfit)
coef(cvfit, s="lambda.min")
coef(cvfit, s="lambda.1se")
```


```
33 x 1 sparse Matrix of class "dgCMatrix"
```


```

              1
(Intercept)  -2.76725511
(Intercept)      .
total_acc      0.05938371
tot_cur_bal    -0.06117577
inq_last_6mths  0.16271536
loan_amnt      0.28691614
dti            -0.04946683
int_rate       0.55755441
total_pymnt    -0.54998352
total_rec_int_log  0.82338753
log_last_pymnt_amnt -0.87092358
log_annual_inc -0.11254141
open_acc       -0.06106200
log_tot_coll_amt -0.10893522
term 60 months -0.67964146
home_ownershipOTHER  0.68068415
home_ownershipOWN    -0.05372664
home_ownershipRENT   0.06047789
application_typeJOINT -0.83170926
state_mean_intlow    -0.07241401
state_mean_intlowmedium  0.01656687
state_mean_intmediumhigh .
emp_length1 year      .
emp_length10+ years   -0.09469249
emp_length2 years     .
```

```
emp_length3 years      .
emp_length4 years      .
emp_length5 years      0.04940558
emp_length6 years      0.13391433
emp_length7 years      0.03648551
emp_length8 years      .
emp_length9 years      .
emp_lengthn/a          -0.04801034

33 x 1 sparse Matrix of class "dgCMatrix"
                                1
(Intercept)                    -2.718571527
(Intercept)                     .
total_acc                      .
tot_cur_bal                    -0.034258341
inq_last_6mths                 0.135462977
loan_amnt                      0.102038385
dti                            .
int_rate                       0.529893971
total_pymnt                    -0.331254575
total_rec_int_log              0.613571692
log_last_pymnt_amnt           -0.686113938
log_annual_inc                 -0.040403510
open_acc                       -0.008251434
log_tot_coll_amt               -0.069412048
term 60 months                 -0.395110780
home_ownershipOTHER            .
home_ownershipOWN              .
home_ownershipRENT             0.045449241
application_typeJOINT          .
state_mean_intlow              .
state_mean_intlowmedium        .
state_mean_intmediumhigh       .
emp_length1 year               .
emp_length10+ years            -0.052798159
emp_length2 years              .
emp_length3 years              .
emp_length4 years              .
emp_length5 years              .
emp_length6 years              0.003031634
emp_length7 years              .
emp_length8 years              .
emp_length9 years              .
emp_lengthn/a                  .
```





In [21]:

```
# apply the model with cvfit$lambda.min to test data
ind_test <- test.sub[, -1]
x_test <- model.matrix(~.,ind_test)
dep_test <- test.sub[, 1]
y_test <- as.matrix(dep_test)
#predict class, type="class"
lasso_prob <- predict(cvfit,newx = x_test,s=cvfit$lambda.min,type="response")
#translate probabilities to predictions
lasso_predict <- rep("neg",nrow(ind_test))
lasso_predict[lasso_prob>.5] <- "pos"
#confusion matrix
table(pred=lasso_predict,true=test.sub$loan_status_binary)
```

	true	
pred	0	1
neg	245732	20124
pos	214	144

In [ ]:

In [ ]: